

Delving Deeper into the Hadoop API

Chapter 4.1



Delving Deeper into the Hadoop API

- Using the ToolRunner class
- Decreasing the amount of intermediate data with Combiners
- Setting up and tearing down Mappers and Reducers using the setup and cleanup methods
- How to access HDFS programmatically
- How to use the distributed cache
- How to use the Hadoop API's library of Mappers, Reducers, and Partitioners

Chapter Topics

Delving Deeper into the Hadoop API

- Using the ToolRunner Class
- Setting Up and Tearing Down Mappers and Reducers
- Decreasing the Amount of Intermediate Data with Combiners
- Accessing HDFS programmatically
- Using the Distributed Cache
- Using the Hadoop API's Library of Mappers, Reducers and Partitioners

Why Use ToolRunner?

- **You can use ToolRunner in MapReduce driver classes**
 - This is not required, but is a best practice
- **ToolRunner uses the GenericOptionsParser class internally**
 - Allows you to specify configuration options on the command line
 - Also allows you to specify items for the Distributed Cache on the command line (see later)

How to Implement ToolRunner: Complete Driver

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.Job; import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WordCount extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new Configuration(), new WordCount(), args);
        System.exit(exitCode);
    }

    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.printf(
                "Usage: %s [generic options] <input dir> <output dir>\n",
                getClass().getSimpleName());
            return -1;
        }
        Job job = new Job(getConf());
        job.setJarByClass(WordCount.class);
        job.setJobName("Word Count");

        FileInputFormat.setInputPaths(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));

        job.setMapperClass(WordMapper.class);
        job.setReducerClass(SumReducer.class);
        job.setMapOutputKeyClass(Text.class);
        job.setMapOutputValueClass(IntWritable.class);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);

        boolean success = job.waitForCompletion(true);
        return success ? 0 : 1;
    }
}
```

How to Implement ToolRunner: Imports

```
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class WordCount {
    public static void main(String[] args) {
        int exitCode = ToolRunner.run(new Configuration(), new WordCount(),
            args);
        System.exit(exitCode);
    }

    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            System.out.printf(
                "Usage: %s [generic options] <input dir> <output dir>\n",
                getClass().getSimpleName());
            return -1;
        }
        Job job = new Job(getConf());
    }
}
```

Import the relevant classes. We omit the import statements in future slides for brevity.

How to Implement ToolRunner: Driver Class Definition

```
public class WordCount extends Configured implements Tool {  
  
    public static void main(String[] args) {  
        int exitCode = ToolRunner.run(getConf(), WordCount.class, args);  
        System.exit(exitCode);  
    }  
  
    public int run(String[] args) throws IOException, InterruptedException,  
        ClassNotFoundException {  
        if (args.length != 2, !args[0].startsWith("-")) {  
            System.out.printf(  
                "Usage: %s [generic options] <input dir> <output dir>\n", getClass().getSimpleName());  
            return -1;  
        }  
        Job job = new Job(getConf());  
        job.setJarByClass(WordCount.class); job.setJobName("Word Count");  
  
        FileInputFormat.setInputPaths(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(SumReducer.class);  
        job.setMapOutputKeyClass(Text.class);  
        job.setMapOutputValueClass(IntWritable.class);  
  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
  
        boolean success = job.waitForCompletion(true);  
        return success ? 0 : 1;  
    }  
}
```

The driver class implements the Tool interface and extends the Configured class.

How to Implement ToolRunner: Main Method

```
public class WordCount extends Configured implements Tool {  
  
    public static void main(String[] args) throws Exception {  
        int exitCode = ToolRunner.run(new Configuration(),  
            new WordCount(), args);  
        System.exit(exitCode);  
    }  
  
    public int run(String[] args) throws IOException, InterruptedException {  
        if (args.length != 2)  
            System.out.printf("Usage: %s [generic options] <input> <output> [m_r, getConf(), getJarByClass(), getMapperClass(), getReducerClass(), getMapOutputKeyClass(), getMapOutputValueClass(), setInputPaths(), setOutputPath(), setMapperClass(), setReducerClass(), setMapOutputKeyClass(), setMapOutputValueClass(), setOutputKeyClass(), setOutputValueClass()]\n", getClass().getName());  
        return -1;  
    }  
    Job job = new Job(getConf());  
    job.setJarByClass(WordCount.class);  
    job.setJobName("Word Count");  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    job.setMapperClass(WordMapper.class);  
    job.setReducerClass(SumReducer.class);  
    job.setMapOutputKeyClass(Text.class);  
    job.setMapOutputValueClass(IntWritable.class);  
  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
  
    ...  
}
```

Stanford CS246H
Winter '15

© Copyright 2010-2014 Cloudera. All rights reserved. Not to be reproduced without prior written consent.

How to Implement ToolRunner: Run Method

```
public class WordCount {  
    public static void main(String[] args) {  
        int exitCode = ToolRunner.run(new WordCount(), args);  
        System.exit(exitCode);  
    }  
}
```

The driver `run` method creates, configures, and submits the job.

```
public int run(String[] args) throws Exception {  
    if (args.length != 2) {  
        System.out.printf(  
            "Usage: %s [generic options] <input dir> <output dir>\n", getClass().getSimpleName());  
        return -1;  
    }  
    Job job = new Job(getConf());  
    job.setJarByClass(WordCount.class);  
    job.setJobName("Word Count");  
  
    FileInputFormat.setInputPaths(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
  
    ...  
}
```

ToolRunner Command Line Options

- ToolRunner allows the user to specify configuration options on the command line
- Commonly used to specify Hadoop properties using the **-D** flag
 - Will override any default or site properties in the configuration
 - But will *not* override those set in the driver code

```
$ hadoop jar myjar.jar MyDriver \
-D mapreduce.job.reduces=10 myinputdir myoutputdir
```

- Note that **-D** options must appear before any additional program arguments
- Can specify an XML configuration file with **-conf**
- Can specify the default filesystem with **-fs uri**
 - Shortcut for **-D fs.defaultFS=uri**

Chapter Topics

Delving Deeper into the Hadoop API

- Using the ToolRunner Class
- **Setting Up and Tearing Down Mappers and Reducers**
- Decreasing the amount of intermediate data with Combiners
- Accessing HDFS programmatically
- Using the Distributed Cache
- Using the Hadoop API's Library of Mappers, Reducers and Partitioners

The setup Method

- It is common to want your Mapper or Reducer to execute some code before the `map` or `reduce` method is called for the first time
 - Initialize data structures
 - Read data from an external file
 - Set parameters
- The `setup` method is run before the `map` or `reduce` method is called for the first time
 -

```
public void setup(Context context)
```

The cleanup Method

- Similarly, you may wish to perform some action(s) after all the records have been processed by your Mapper or Reducer
- The `cleanup` method is called before the Mapper or Reducer terminates

```
public void cleanup(Context context) throws  
    IOException, InterruptedException
```

Passing Parameters

```
public class MyDriverClass {  
    public int main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        conf.setInt ("paramname", value);  
        Job job = new Job(conf);  
        ...  
        boolean success = job.waitForCompletion(true);  
        return success ? 0 : 1;  
    }  
}
```

```
public class MyMapper extends Mapper {  
  
    public void setup(Context context) {  
        Configuration conf = context.getConfiguration();  
        int myParam = conf.getInt("paramname", 0);  
        ...  
    }  
    public void map...  
}
```

Chapter Topics

Delving Deeper into the Hadoop API

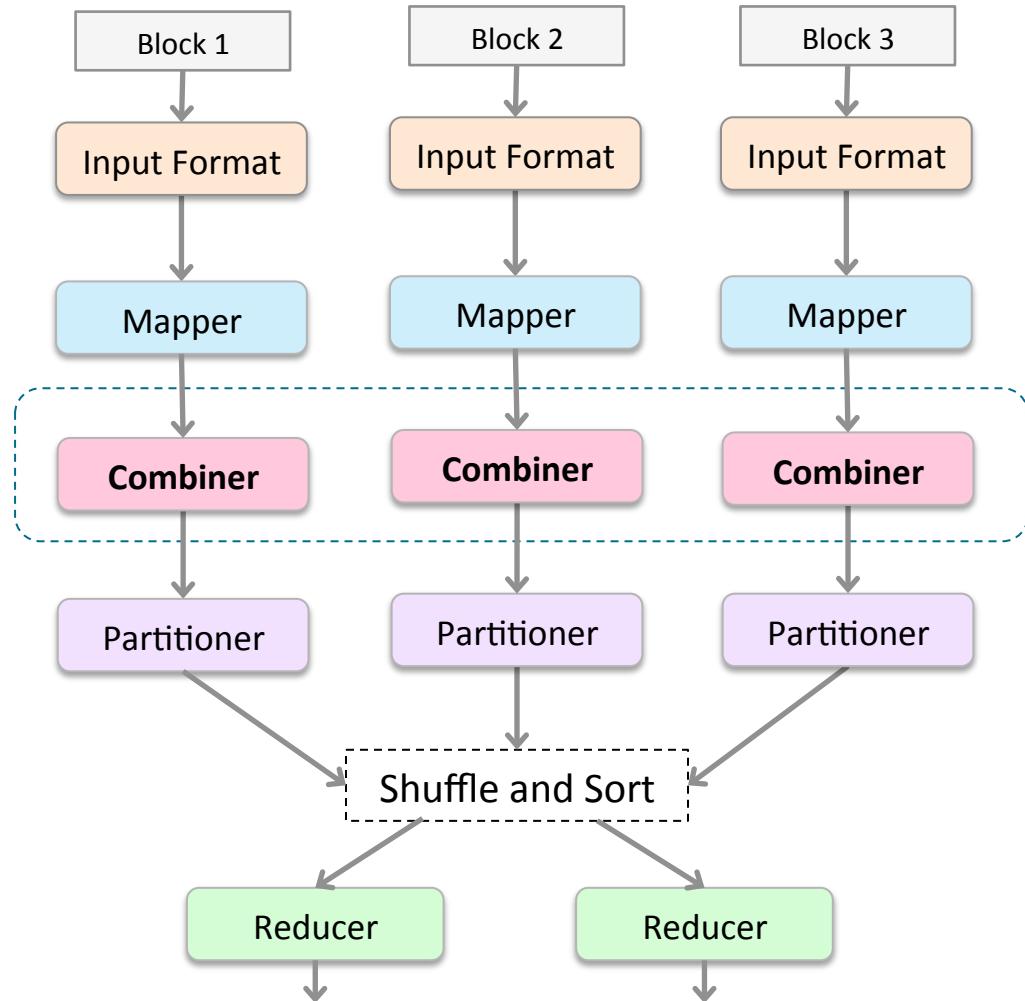
- Using the ToolRunner Class
- Setting Up and Tearing Down Mappers and Reducers
- **Decreasing the Amount of Intermediate Data with Combiners**
- Accessing HDFS programmatically
- Using the Distributed Cache
- Using the Hadoop API's Library of Mappers, Reducers and Partitioners

The Combiner

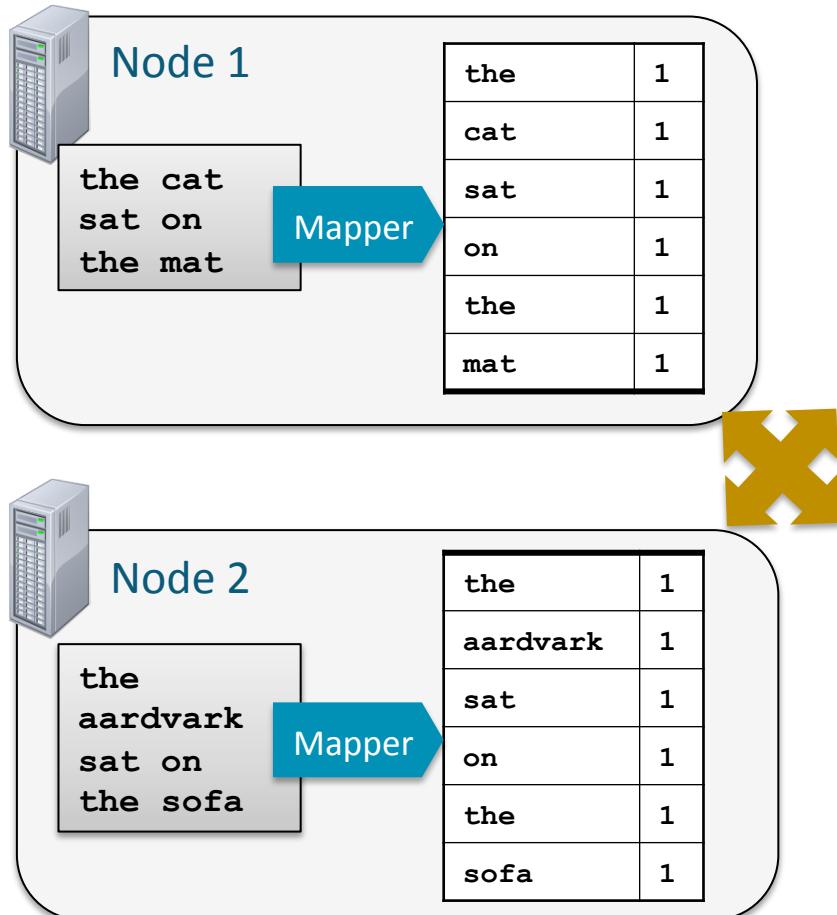
- Often, Mappers produce large amounts of intermediate data
 - That data must be passed to the Reducers
 - This can result in a lot of network traffic
- It is often possible to specify a *Combiner*
 - Like a ‘mini-Reducer’
 - Runs locally on a single Mapper’s output
 - Output from the Combiner is sent to the Reducers
- Combiner and Reducer code are often identical
 - Technically, this is possible if the operation performed is commutative and associative
 - Input and output data types for the Combiner/Reducer must be identical

The Combiner

- Combiners run as part of the Map phase
- Output from the Combiners is passed to the Reducers



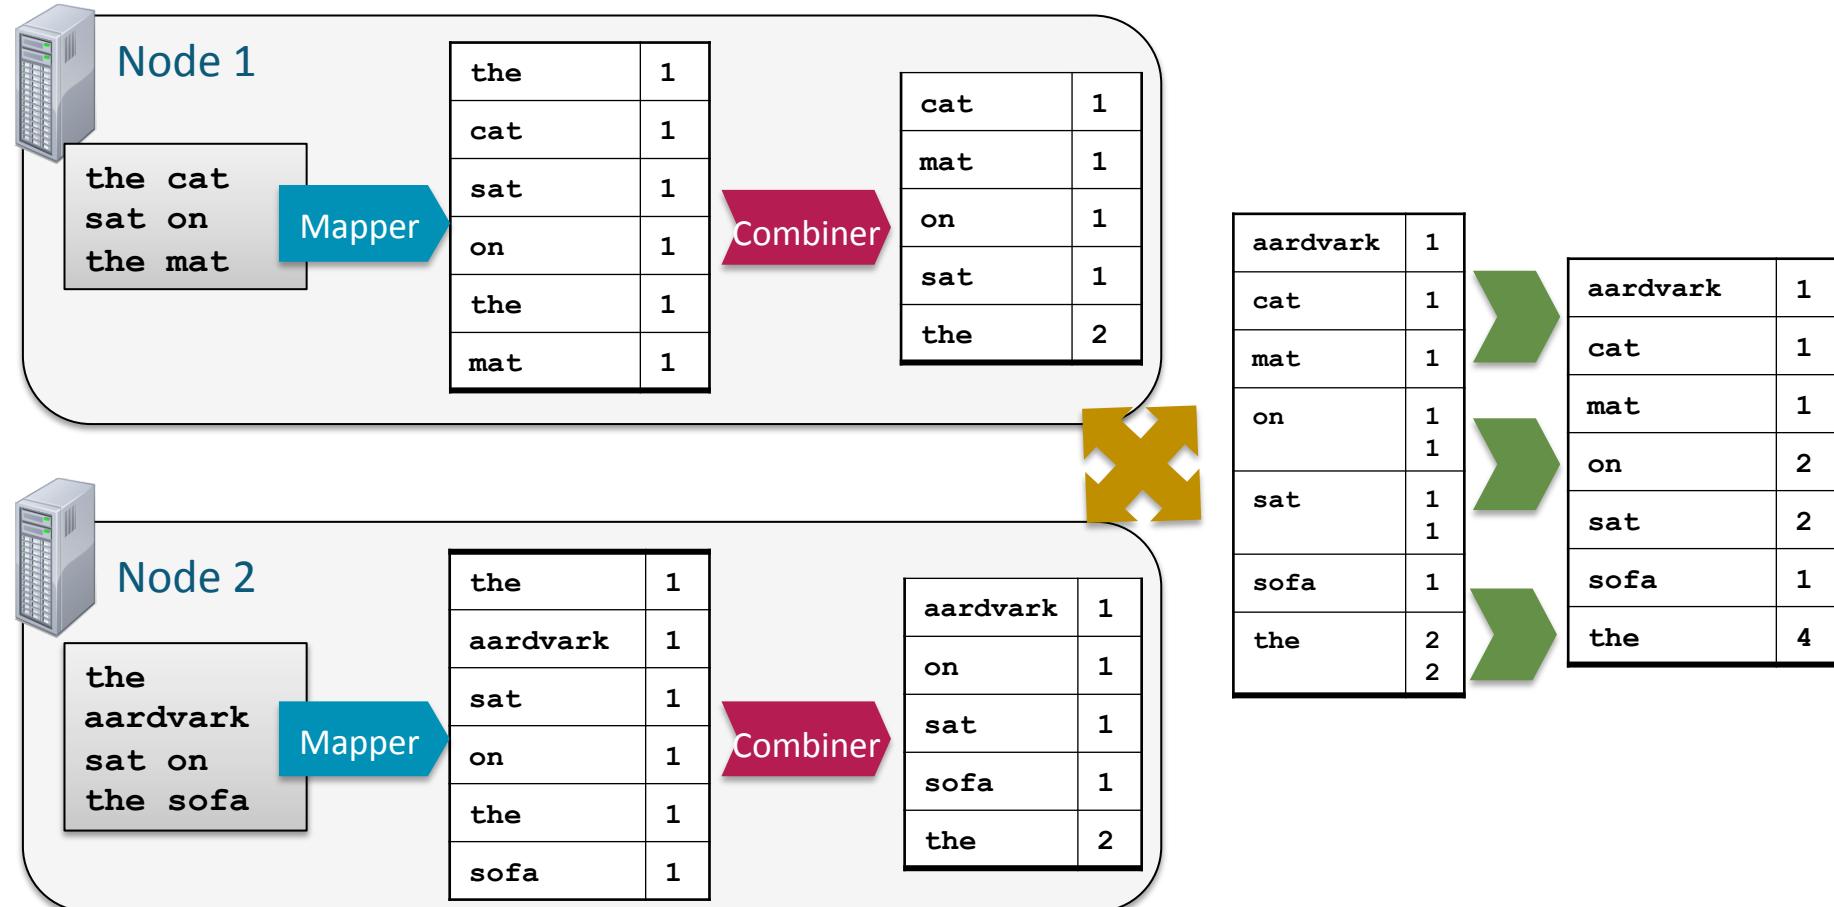
WordCount Revisited



aardvark	1
cat	1
mat	1
on	1
sat	1



WordCount With Combiner



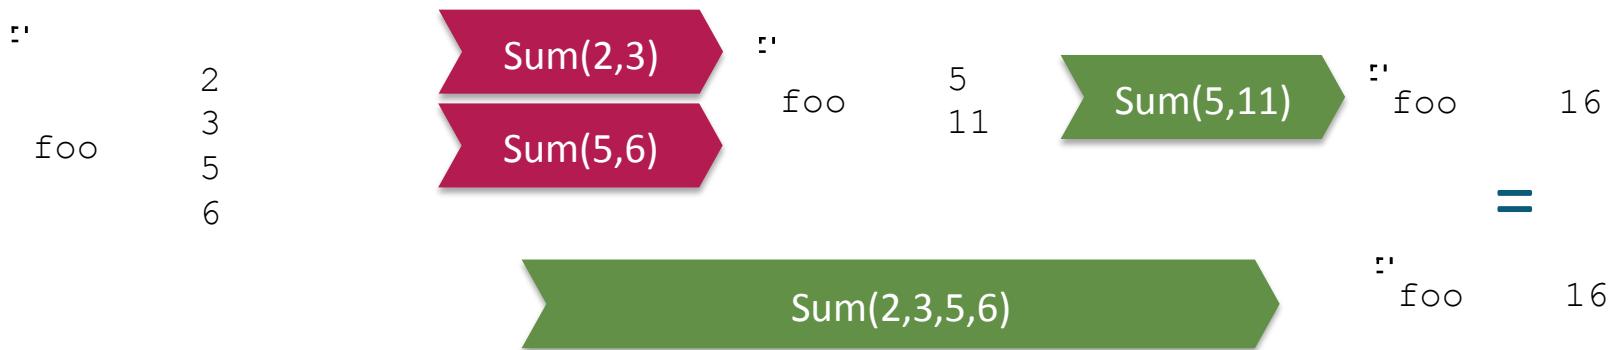
Writing a Combiner

- **The Combiner uses the same signature as the Reducer**
 - Takes in a key and a list of values
 - Outputs zero or more (key, value) pairs
 - The actual method called is the `reduce` method in the class

```
reduce(inter_key, [v1, v2, ...]) →  
      (result_key, result_value)
```

Combiners and Reducers

- Some Reducers may be used as Combiners
 - If operation is associative and commutative, e.g., SumReducer



- Some Reducers *cannot* be used as a Combiner, e.g., AverageReducer



Specifying a Combiner

- Specify the Combiner class to be used in your MapReduce code in the driver
 - Use the `setCombinerClass` method, e.g.:

```
job.setMapperClass(WordMapper.class);  
job.setReducerClass(SumReducer.class);  
job.setCombinerClass(SumReducer.class);
```

- Input and output data types for the Combiner and the Reducer for a job must be identical
- VERY IMPORTANT: The Combiner may run once, or more than once, on the output from any given Mapper
 - Do not put code in the Combiner which could influence your results if it runs more than once

Chapter Topics

Delving Deeper into the Hadoop API

- Using the ToolRunner Class
- Setting Up and Tearing Down Mappers and Reducers
- Decreasing the Amount of Intermediate Data with Combiners
- **Accessing HDFS Programmatically**
- Using the Distributed Cache
- Using the Hadoop API's Library of Mappers, Reducers and Partitioners

Accessing HDFS Programmatically

- **In addition to using the command-line shell, you can access HDFS programmatically**
 - Useful if your code needs to read or write ‘side data’ in addition to the standard MapReduce inputs and outputs
 - Or for programs outside of Hadoop which need to read the results of MapReduce jobs
- **Beware: HDFS is not a general-purpose filesystem!**
 - Files cannot be modified once they have been written, for example
- **Hadoop provides the `FileSystem` abstract base class**
 - Provides an API to generic file systems
 - Could be HDFS
 - Could be your local file system
 - Could even be, for example, Amazon S3

The FileSystem API (1)

- In order to use the `FileSystem` API, retrieve an instance of it

```
Configuration conf = new Configuration();
FileSystem fs = FileSystem.get(conf);
```

- The `conf` object has read in the Hadoop configuration files, and therefore knows the address of the NameNode
- A file in HDFS is represented by a `Path` object

```
Path p = new Path("/path/to/my/file");
```

The FileSystem API (2)

- Some useful API methods:

- `FSDataOutputStream create(...)`
 - Extends `java.io.DataOutputStream`
 - Provides methods for writing primitives, raw bytes etc
- `FSDataInputStream open(...)`
 - Extends `java.io.DataInputStream`
 - Provides methods for reading primitives, raw bytes etc
- `boolean delete(...)`
- `boolean mkdirs(...)`
- `void copyFromLocalFile(...)`
- `void copyToLocalFile(...)`
- `FileStatus[] listStatus(...)`

The FileSystem API: Directory Listing

- Get a directory listing:

```
Path p = new Path("/my/path");

Configuration conf = new Configuration();
FileSystem fs = FileSystem.get(conf);
FileStatus[] fileStats = fs.listStatus(p);

for (int i = 0; i < fileStats.length; i++) {
    Path f = fileStats[i].getPath();

    // do something interesting
}
```

The FileSystem API: Writing Data

- Write data to a file

```
Configuration conf = new Configuration();
FileSystem fs = FileSystem.get(conf);

Path p = new Path("/my/path/foo");

FSDataOutputStream out = fs.create(p, false);

// write some raw bytes
out.writegetBytes();

// write an int
out.writeInt(getInt());

...

out.close();
```

Chapter Topics

Delving Deeper into the Hadoop API

- Using the ToolRunner Class
- Setting Up and Tearing Down Mappers and Reducers
- Decreasing the Amount of Intermediate Data with Combiners
- Accessing HDFS Programmatically
- **Using the Distributed Cache**
- Using the Hadoop API's Library of Mappers, Reducers and Partitioners

The Distributed Cache: Motivation

- A common requirement is for a Mapper or Reducer to need access to some ‘side data’
 - Lookup tables
 - Dictionaries
 - Standard configuration values
- One option: read directly from HDFS in the `setup` method
 - Using the API seen in the previous section
 - Works, but is not scalable
- The Distributed Cache provides an API to push data to all slave nodes
 - Transfer happens behind the scenes before any task is executed
 - Data is only transferred once to each node
 - Note: Distributed Cache is read-only
 - Files in the Distributed Cache are automatically deleted from slave nodes when the job finishes

Using the Distributed Cache: The Difficult Way

- Place the files into HDFS
- Configure the Distributed Cache in your driver code

```
Configuration conf = new Configuration();
DistributedCache.addCacheFile(new URI("/myapp/lookup.dat"),conf);
DistributedCache.addFileToClassPath(new Path("/myapp/mylib.jar"),conf);
DistributedCache.addCacheArchive(new URI("/myapp/map.zip"),conf));
DistributedCache.addCacheArchive(new URI("/myapp/mytar.tar"),conf));
DistributedCache.addCacheArchive(new URI("/myapp/mytgz.tgz"),conf));
DistributedCache.addCacheArchive(new URI("/myapp/mytargz.tar.gz"),conf));
```

- .jar files added with addFileToClassPath will be added to your Mapper or Reducer's classpath
- Files added with addCacheArchive will automatically be dearchived/decompressed

Using the DistributedCache: The Easy Way

- If you are using ToolRunner, you can add files to the Distributed Cache directly from the command line when you run the job
 - No need to copy the files to HDFS first
- Use the **-files** option to add files

```
hadoop jar myjar.jar MyDriver -files file1, file2, file3, ...
```

- The **-archives** flag adds archived files, and automatically unarchives them on the destination machines
- The **-libjars** flag adds jar files to the classpath

Accessing Files in the Distributed Cache

- **Files added to the Distributed Cache are made available in your task's local working directory**
 - Access them from your Mapper or Reducer the way you would read any ordinary local file

```
File f = new File("file_name_here");
```

Chapter Topics

Delving Deeper into the Hadoop API

- Using the ToolRunner Class
- Setting Up and Tearing Down Mappers and Reducers
- Decreasing the Amount of Intermediate Data with Combiners
- Accessing HDFS Programmatically
- Using the Distributed Cache
- **Using the Hadoop API's Library of Mappers, Reducers and Partitioners**

Reusable Classes for the New API

- The `org.apache.hadoop.mapreduce.lib.*/*` packages contain a library of Mappers, Reducers, and Partitioners supporting the new API
- Example classes:
 - `InverseMapper` – Swaps keys and values
 - `RegexMapper` – Extracts text based on a regular expression
 - `IntSumReducer`, `LongSumReducer` – Add up all values for a key
 - `TotalOrderPartitioner` – Reads a previously-created partition file and partitions based on the data from that file
 - Sample the data first to create the partition file
 - Allows you to partition your data into n partitions without hard-coding the partitioning information
- Refer to the Javadoc for classes available in your version of CDH
 - Available classes vary greatly from version to version

Key Points

- **Use the ToolRunner class to build drivers**
 - Parses job options and configuration variables automatically
- **Override Mapper and Reducer setup and cleanup methods**
 - Set up and tear down, e.g. reading configuration parameters
- **Combiners are ‘mini-reducers’**
 - Run locally on Mapper output to reduce data sent to Reducers
- **The FileSystem API lets you read and write HDFS files programmatically**
- **The Distributed Cache lets you copy local files to worker nodes**
 - Mappers and Reducers can access directly as regular files
- **Hadoop includes a library of predefined Mappers, Reducers, and Partitioners**

Bibliography

The following offer more information on topics discussed in this chapter

- Combiners are discussed in TDG 3e on pages 33-36.
- A table describing available filesystems in Hadoop is on pages 52-53 of TDG 3e.
- The HDFS API is described in TDG 3e on pages 55-67.
- Distributed cache: See pages 289-295 of TDG 3e for more details.

Practical Development Tips and Techniques

Chapter 5.1



Practical Development Tips and Techniques

- Strategies for debugging MapReduce code
- How to test MapReduce code locally using LocalJobRunner
- How to write and view log files
- How to retrieve job information with counters
- Why reusing objects is a best practice
- How to create Map-only MapReduce jobs

Chapter Topics

Practical Development Tips and Techniques

- **Strategies for Debugging MapReduce Code**
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs

Introduction to Debugging

- **Debugging MapReduce code is difficult!**
 - Each instance of a Mapper runs as a separate task
 - Often on a different machine
 - Difficult to attach a debugger to the process
 - Difficult to catch ‘edge cases’
- **Very large volumes of data mean that unexpected input is likely to appear**
 - Code which expects all data to be well-formed is likely to fail

Common-Sense Debugging Tips

- **Code defensively**
 - Ensure that input data is in the expected format
 - Expect things to go wrong
 - Catch exceptions
- **Start small, build incrementally**
- **Make as much of your code as possible Hadoop-agnostic**
 - Makes it easier to test
- **Write unit tests**
- **Test locally whenever possible**
 - With small amounts of data
- **Then test in pseudo-distributed mode**
- **Finally, test on the cluster**

Testing Strategies

- When testing in pseudo-distributed mode, ensure that you are testing with a similar environment to that on the real cluster
 - Same amount of RAM allocated to the task JVMs
 - Same version of Hadoop
 - Same version of Java
 - Same versions of third-party libraries

Chapter Topics

Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- **Testing MapReduce Code Locally Using LocalJobRunner**
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs

Testing Locally (1)

- **Hadoop can run MapReduce in a single, local process**
 - Does not require any Hadoop daemons to be running
 - Uses the local filesystem instead of HDFS
 - Known as LocalJobRunner mode
- **This is a very useful way of quickly testing incremental changes to code**

Testing Locally (2)

- To run in LocalJobRunner mode, add the following lines to the driver code:

```
Configuration conf = new Configuration();
conf.set("mapreduce.framework.name", "local");
conf.set("fs.defaultFS", "file:///");
```

- Or set these options on the command line if your driver uses ToolRunner
 - fs** is equivalent to **-D fs.defaultFS**
 - jt** is equivalent to **-D mapreduce.framework.name**
 - e.g.

```
$ hadoop jar myjar.jar MyDriver -fs=file:/// -jt=local \
indir outdir
```

-fs seems flakey in YARN, so best to prepend paths with file://

Testing Locally (3)

- **Some limitations of LocalJobRunner mode:**
 - Distributed Cache does not work
 - The job can only specify a single Reducer
 - Some ‘beginner’ mistakes may not be caught
 - For example, attempting to share data between Mappers will work, because the code is running in a single JVM

Chapter Topics

Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- **Writing and Viewing Log Files**
- Retrieving Job Information with Counters
- Reusing Objects
- Creating Map-only MapReduce Jobs

Before Logging: `stdout` and `stderr`

- Tried-and-true debugging technique: write to `stdout` or `stderr`
- If running in LocalJobRunner mode, you will see the results of `System.err.println()`
- If running on a cluster, that output will not appear on your console
 - Output is visible via Hadoop's Web UI

Aside: The Hadoop Web UI

- All Hadoop daemons contain a Web server
 - Exposes information on a well-known port
- Most important for developers is the k U Web UI
 - `http://<master_address>:50070/`
 - `http://localhost:50070/` if running in pseudo-distributed mode
- Also useful: the NameNode Web UI
 - `http://<name_node_address>:50070/`

Aside: The Hadoop Web UI (cont'd)

- Your instructor will now demonstrate the ResourceManager UI

The screenshot shows a Mozilla Firefox browser window displaying the "Attempts for task_1420746143356_0003_m_000000" page. The URL is http://quickstart.cloudera:19888/jobhistory/task/task_1420746143356_0003_m_000000. The page title is "Attempts for task_1420746143356_0003_m_000000". The left sidebar has a tree view with "Application", "Job", "Task" (selected), and "Tools". The main content area shows a table with one entry:

Attempt	State	Status	Node	Logs	Start Time	Finish Time	Elapsed Time	Note
attempt_1420746143356_0003_m_000000_0	SUCCEEDED	Records R/W=15/1	default-rack/quickstart.cloudera:8042	logs	Fri Jan 16 11:56:50 2015	Fri Jan 16 11:57:10 2015	<000 2015	

A red circle highlights the "Logs" column header. The bottom of the page includes links for "About Apache Hadoop" and the terminal prompt "[cloudera@quickstart:~]".

Logging: Better Than Printing

- **println statements rapidly become awkward**
 - Turning them on and off in your code is tedious, and leads to errors
- **Logging provides much finer-grained control over:**
 - What gets logged
 - When something gets logged
 - How something is logged



Logging With log4j

- Hadoop uses log4j to generate all its log files
- Your Mappers and Reducers can also use log4j
 - All the initialization is handled for you by Hadoop
- Add the log4j.jar-<version> file from your CDH distribution to your classpath when you reference the log4j classes

```
import org.apache.log4j.Level;
import org.apache.log4j.Logger;

class FooMapper implements Mapper {
    private static final Logger LOGGER =
        Logger.getLogger (FooMapper.class.getName());
    ...
}
```

Logging With log4j (cont'd)

- Simply send strings to loggers tagged with severity levels:

```
LOGGER.trace("message");
LOGGER.debug("message");
LOGGER.info("message");
LOGGER.warn("message");
LOGGER.error("message");
```

- Beware expensive operations like concatenation

- To avoid performance penalty, make it conditional like this:

```
if (LOGGER.isDebugEnabled()) {
    LOGGER.debug("Account info:" + acct.getReport());
}
```

log4j Configuration

- Node-wide configuration for log4j is stored in `/etc/hadoop/conf/log4j.properties`
- Override settings for your application in your own `log4j.properties`
 - Can change global log settings with `hadoop.root.log` property
 - Can override log level on a per-class basis, e.g.

```
log4j.logger.org.apache.hadoop.mapred.JobTracker=WARN
```

```
log4j.logger.com.mycompany.myproject.FooMapper=DEBUG
```



Full class name

- Or set the level programmatically:

```
LOGGER.setLevel(Level.WARN);
```

Setting Logging Levels for a Job

- You can tell Hadoop to set logging levels for a job using configuration properties

- `-mapred.map.child.log.level`
 - `-mapred.reduce.child.log.level`

- Examples

- Set the logging level to DEBUG for the Mapper

```
$ hadoop jar myjob.jar MyDriver \
-Dmapred.map.child.log.level=DEBUG indir outdir
```

- Set the logging level to WARN for the Reducer

```
$ hadoop jar myjob.jar MyDriver \
-Dmapred.reduce.child.log.level=WARN indir outdir
```

Where Are Log Files Stored?

- **Log files are stored on the machine where the task attempt ran**
 - Location is configurable
 - By default:
`/var/log/hadoop-0.20-mapreduce/
userlogs/${task.id}/syslog`
- **You will often not have ssh access to a node to view its logs**
 - Much easier to use the ResourceManager Web UI
 - Automatically retrieves and displays the log files for you

Restricting Log Output

- If you suspect the input data of being faulty, you may be tempted to log the (key, value) pairs your Mapper receives
 - Reasonable for small amounts of input data
 - Caution! If your job runs across 500GB of input data, you could be writing up to 500GB of log files!
 - Remember to think at scale...
- Instead, wrap vulnerable sections of code in `try { . . . }` blocks
 - Write logs in the `catch { . . . }` block
 - This way only critical data is logged



Aside: Throwing Exceptions

- You could throw exceptions if a particular condition is met
 - For example, if illegal data is found

```
throw new RuntimeException("Your message here");
```



- Usually not a good idea
 - Exception causes the task to fail
 - If a task fails four times, the entire job will fail

Chapter Topics

Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- **Retrieving Job Information with Counters**
- Reusing Objects
- Creating Map-only MapReduce Jobs

What Are Counters? (1)

- **Counters provide a way for Mappers or Reducers to pass aggregate values back to the driver after the job has completed**
 - Their values are also visible from the JobTracker's Web UI
 - And are reported on the console when the job ends
- **Very basic: just have a name and a value**
 - Value can be incremented within the code
- **Counters are collected into Groups**
 - Within the group, each Counter has a name
- **Example: A group of Counters called RecordType**
 - Names: TypeA, TypeB, TypeC
 - Appropriate Counter can be incremented as each record is read in the Mapper

What Are Counters? (2)

- Counters can be set and incremented via the method

```
context.getCounter(group, name).increment(amount);
```

- Example:

```
context.getCounter("RecordType", "A").increment(1);
```

Retrieving Counters in the Driver Code

- To retrieve Counters in the Driver code after the job is complete, use code like this in the driver:

```
long typeARecords =  
    job.getCounters().findCounter("RecordType", "A").getValue();  
  
long typeBRecords =  
    job.getCounters().findCounter("RecordType", "B").getValue();
```

Counters: Caution

- **Do not rely on a counter's value from the Web UI while a job is running**
 - Due to possible speculative execution, a counter's value could appear larger than the actual final value
 - Modifications to counters from subsequently killed/failed tasks will be removed from the final count

Chapter Topics

Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- **Reusing Objects**
- Creating Map-only MapReduce Jobs

Reuse of Objects is Good Practice (1)

- It is generally good practice to reuse objects
 - Instead of creating many new objects
- Example: Our original WordCount Mapper code

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        String word = value.toString();
        for (String line : lines) {
            if (line.startsWith(word)) {
                context.write(new Text(word), new IntWritable(1));
            }
        }
    }
}
```

Each time the map () method is called, we create a new Text object and a new IntWritable object.

Reuse of Objects is Good Practice (2)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        String line = value.toString();

        for (String word : line.split("\\w+")) {
            if (word.length() > 0) {
                wordObject.set(word);
                context.write(wordObject, one);
            }
        }
    }
}
```

Create objects for the key and value outside of your map () method

Reuse of Objects is Good Practice (3)

- Instead, this is better practice:

```
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static IntWritable one = new IntWritable(1);
    private Text wordObject = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {
        if (word.length() > 0) {
            wordObject.set(word);
            context.write(wordObject, one);
        }
    }
}
```

Within the `map()` method, populate the objects and write them out. Hadoop will take care of serializing the data so it is perfectly safe to re-use the objects.

Object Reuse: Caution!

- Hadoop re-uses objects all the time
- For example, each time the Reducer is passed a new value, the same object is reused
- This can cause subtle bugs in your code
 - For example, if you build a list of value objects in the Reducer, each element of the list will point to the same underlying object
 - Unless you do a deep copy

Chapter Topics

Practical Development Tips and Techniques

- Strategies for Debugging MapReduce Code
- Testing MapReduce Code Locally Using LocalJobRunner
- Writing and Viewing Log Files
- Retrieving Job Information with Counters
- Reusing Objects
- **Creating Map-only MapReduce jobs**

Map-Only MapReduce Jobs

- There are many types of job where only a Mapper is needed
- Examples:
 - Image processing
 - File format conversion
 - Input data sampling
 - ETL

Creating Map-Only Jobs

- To create a Map-only job, set the number of Reducers to 0 in your Driver code

```
job.setNumReduceTasks(0);
```

- Call the Job.setOutputKeyClass and Job.setOutputValueClass methods to specify the output types
 - Not the Job.setMapOutputKeyClass and Job.setMapOutputValueClass methods
- Anything written using the Context.write method in the Mapper will be written to HDFS
 - Rather than written as intermediate data
 - One file per Mapper will be written

Key Points

- LocalJobRunner lets you test jobs on your local machine
- Hadoop uses the Log4J framework for logging
- Reusing objects is a best practice
- Counters provide a way of passing numeric data back to the driver
- Create Map-only MapReduce jobs by setting the number of Reducers to zero

Bibliography

The following offer more information on topics discussed in this chapter

- Java version selection

- <http://wiki.apache.org/hadoop/HadoopJavaVersions>

- For an example of image processing and file format conversion in Hadoop, see

- <http://open.blogs.nytimes.com/2007/11/01/self-service-prorated-super-computing-fun/>